

Benchmark Averages Hide the Failures That Matter: Quantizing ESM-2 for Protein Variant-Effect Prediction

Qing Shao

Department of Chemical and Materials Engineering
University of Kentucky, Lexington, KY 40506

August 9, 2026

Abstract

We benchmark six numerical precision configurations for ESM-2 protein language models across three axes — throughput, memory footprint, and predictive accuracy — on two workloads with sharply different characteristics: bulk embedding extraction and deep mutational scanning (DMS) variant-effect scoring. For ESM2-3B on an H200, FP8 dynamic quantization (`fp8_dyn`) delivers $1.15\times$ the throughput of bf16 while halving peak memory ($12.63 \rightarrow 7.24$ GB), making it the best configuration for embedding extraction. The same configuration is a poor choice for DMS scoring, where bf16 in eager mode is fastest at every problem size and quantization is pure overhead. Accuracy is evaluated on the complete ProteinGym substitution benchmark — 201 assays, 2.41M variants, both model scales — with a paired bootstrap clustered on protein. Two findings follow. First, fidelity measured *against fp32*, the metric such studies normally report, predicts the *magnitude* of the change in real-world accuracy ($r = 0.81$ over 1005 assay/configuration pairs) but not its *sign*: INT4 weight-only quantization is significantly *beneficial* at 3B and significantly *harmful* at 650M. Second, and of more practical consequence, no configuration shifts mean benchmark correlation by more than 0.007, yet INT8 dynamic quantization — indistinguishable from fp32 on that mean ($p = 0.34$) — takes a single assay from $\rho = 0.591$ to 0.223. Benchmark averages conceal precisely the failure that determines whether a configuration is deployable. We conclude that drift-from-fp32 metrics are sound risk bounds but unsound quality rankings, that configuration selection should be made on worst-case rather than mean behaviour, and we report four measurement artifacts encountered during this study, three of which inverted the result they were meant to measure.

Keywords: protein language models; ESM-2; post-training quantization; variant-effect prediction; ProteinGym; benchmark methodology; FP8; INT8

1 Introduction

Quantization is usually presented as a straightforward trade: reduce numerical precision, gain speed and memory, lose a little accuracy. This framing is inherited from large autoregressive language models, where it is broadly correct. We set out to test whether it transfers to ESM-2, a family of protein language models, and found that essentially every part of it requires qualification.

ESM-2 is a BERT-style bidirectional *encoder*. It performs a single forward pass with no KV cache and no autoregressive decode. This places it in a **compute-bound** regime rather than the memory-bandwidth-bound regime that governs LLM inference at batch size 1. The distinction is

not academic: the dominant open-source quantization toolchains (GPTQ, AWQ, bitsandbytes NF4, GGUF) are all *weight-only*. They store INT4/INT8 weights and dequantize to bf16 immediately before an ordinary bf16 matrix multiply. When the bottleneck is memory bandwidth, this is a large win. When the bottleneck is arithmetic, as it is here, it adds work to an already-saturated kernel.

This paper addresses three questions:

1. What do the available quantization configurations actually cost and deliver on the three axes of speed, memory, and accuracy?
2. Does the answer depend on the workload? (It does, and the two workloads we study prefer opposite configurations.)
3. Do the fidelity metrics conventionally used to validate quantization predict real-world predictive accuracy? (Only partially, and in a way that matters.)

Contributions

- An evaluation of six precision configurations for ESM-2 on the **complete ProteinGym substitution benchmark** — 201 assays, 2.41M variants, 40,775 masked positions, at two model scales — rather than on drift from the full-precision model or on a handful of assays.
- The finding that **benchmark-mean accuracy is nearly uninformative for deployment**. No configuration shifts mean correlation by more than 0.007, while one configuration that passes the mean test at $p = 0.34$ takes a single assay from $\rho = 0.591$ to 0.223. We argue that worst-case single-assay drift, not the mean, is the selection criterion.
- A quantitative characterisation of what drift-from-fp32 metrics do and do not buy: over 1005 assay/configuration pairs they predict the magnitude of ground-truth change ($r = 0.81$) but not its direction, which does not even hold its sign between model scales.
- The observation that the **two common ESM-2 workloads prefer opposite configurations**, so that a pipeline running both cannot be tuned with a single choice.
- A documented account of **four measurement artifacts**, three of which inverted the quantity being measured while producing plausible, monotonic result tables — offered as reusable hazards rather than as local mistakes.
- Open code, per-variant predictions for all 12 model/configuration combinations, and per-assay statistics (Section 9).

2 Background and Related Work

Protein language models. ESM-2 [7] is a family of transformer encoders trained with masked language modelling on UniRef, following ESM-1b [12]. Architecturally these are BERT-style bidirectional encoders [2]: a single forward pass, no KV cache, no autoregressive decode. Meier et al. [9] established that such models predict variant effects zero-shot via *masked marginals* — masking a mutated position and comparing the log-probabilities of mutant and wild-type residues — which is the scoring protocol used throughout this work.

Variant-effect benchmarks. ProteinGym [11] aggregates deep mutational scanning assays into a standard benchmark; the v1 substitution set used here contains 217 assays spanning 2.47M variants and five functional categories, including the mega-scale stability measurements of Tsuboyama et al. [15]. Reporting mean Spearman correlation across assays is the established summary statistic, and one of our contributions is to show what that statistic conceals.

Post-training quantization. The dominant open-source toolchains are weight-only: GPTQ [6] and AWQ [8] store INT4 weights and dequantize before a bf16 matmul, while LLM.int8() [3] introduced mixed-precision decomposition to handle activation outliers. These were designed for autoregressive decoding, where batch-1 inference is memory-bandwidth-bound and dequantization is effectively free. Weight-activation (W8A8) schemes and the FP8 E4M3/E5M2 formats [10] instead reduce arithmetic cost, which is what a compute-bound encoder requires. We use `torchao` [14] for all configurations and PyTorch 2’s compiler stack [1], whose Inductor backend emits Triton [13] kernels; Section 4.1 shows that the compile mode, not the quantization scheme, determines whether W8A8 is viable at all. Quantization is also load-bearing for parameter-efficient fine-tuning [4], a setting we do not evaluate.

Positioning. Quantization studies for protein language models are comparatively rare, and where accuracy is reported it is typically drift against the full-precision model on a small probe set. Our contribution is to replace that proxy with the full downstream benchmark at two model scales, and to show that the proxy is a valid risk bound but an invalid quality ranking.

3 Experimental Design

3.1 Configurations

Six configurations were evaluated, all applied via `torchao` to the encoder `Linear` layers only:

Table 1: Quantization configurations. Weight-only variants target memory; the dynamic (W8A8) variants also quantize activations and can therefore use low-precision tensor cores for the matrix multiply itself.

Configuration	Class	Description
<code>fp32</code>	baseline	Reference for all fidelity measurements.
<code>bf16</code>	baseline	The practical production default.
<code>int8_wo</code>	weight-only	INT8 weights, bf16 compute. Memory-oriented.
<code>int4_wo</code>	weight-only	INT4 weights, bf16 compute. Smallest footprint.
<code>int8_dyn</code>	W8A8	INT8 weights <i>and</i> dynamically quantized activations; uses INT8 tensor cores.
<code>fp8_dyn</code>	W8A8	FP8 (E4M3) weights and activations. Requires compute capability ≥ 8.9 ; H200 only.

3.2 Workloads

Bulk embedding extraction. Large length-bucketed batches under a token budget, measured in residues per second. This is the compute-bound regime.

DMS variant-effect scoring. The standard ESM masked-marginals protocol: for each mutated position, mask it and record $\log p(\text{mut}) - \log p(\text{wt})$ from the resulting distribution. Cost scales with

the *number of mutated positions*, not with the number of variants, and the masked batch is narrow. This is the latency-bound regime.

Algorithm 1 states the scoring procedure as implemented. Two optimizations, both absent from typical reference implementations, are marked: only positions carrying a variant are masked (line 5), and the masked copies are batched rather than run one at a time (line 6). Together they are worth one to two orders of magnitude on a real assay — the median ProteinGym assay mutates 96 positions out of $L = 222$ — and they are applied identically to every configuration, so they change the absolute cost of the benchmark without affecting any comparison within it.

Algorithm 1 Masked-marginals variant scoring, as benchmarked. Cost is $\lceil |P|/B \rceil$ forward passes and is independent of $|V|$, so an assay with 16,000 variants over 900 positions costs the same as one with 900.

Require: wild-type sequence s , variants V , model f_θ , batch size B

Ensure: score $\sigma(v)$ for every $v \in V$

```

1:  $M \leftarrow \{ \text{parse}(v) : v \in V \}$  ▷ each  $v$  is a set of (wt, pos, mut) triples
2: for each  $(a, i, b) \in \bigcup M$  do
3:   assert  $s_i = a$  ▷ indexing check; a mismatch is silent and unrecoverable later
4: end for
5:  $P \leftarrow \{ i : (a, i, b) \in \bigcup M \}$  ▷ only variant-carrying positions, not all  $L$ 
6: for each chunk  $C \subseteq P$  with  $|C| = B$  do
7:    $X \leftarrow |C|$  copies of  $\text{tokenize}(s)$ ; set  $X_{r, c_r} \leftarrow [\text{MASK}]$  ▷ one masked position per row
8:    $Z \leftarrow f_\theta(X)$  ▷ single batched forward pass
9:   for each  $r$ , with  $c_r \in C$  do
10:     $\ell_{c_r} \leftarrow \log \text{softmax}(Z_{r, c_r} \text{ in fp32})$  ▷ fp32 softmax regardless of compute precision
11:   end for
12: end for
13: for each  $v \in V$  do
14:    $\sigma(v) \leftarrow \sum_{(a, i, b) \in M_v} (\ell_i[b] - \ell_i[a])$  ▷ difference of near-equal terms: noise does not cancel
15: end for
```

The final line is why DMS ρ is the most sensitive fidelity metric we record. The score is a difference of two nearly equal log-probabilities, so quantization error in ℓ_i does not cancel between the two terms; it is amplified relative to the signal.

3.3 Metrics

Three fidelity metrics were recorded, deliberately ordered by sensitivity:

1. **Embedding cosine similarity** (mean-pooled, excluding special tokens and padding). For-giving: error averages out along the sequence.
2. **Logit KL divergence**, per residue. Moderately sensitive.
3. **DMS Spearman ρ against fp32**. Most sensitive, because the score is a *difference of two near-equal log-probabilities*, so quantization noise does not cancel and is amplified relative to signal.

All three are drift-from-fp32 measures. They quantify how far a configuration moves the answer, not whether it moves toward or away from the truth. To measure the latter, we added a fourth

metric requiring real experimental data: **Spearman ρ against wet-lab DMS measurements**, from ProteinGym.

3.4 Real assay data

Three ProteinGym substitution assays were selected to span the sequence-length range, since peak memory during masked-marginals scoring grows with L :

Table 2: ProteinGym assays used. All 19,557 single-substitution variants had their stated wild-type residue verified against the reference `target_seq` before use; all matched. This check is load-bearing — an off-by-one in position indexing does not raise an error, it silently scores the wrong positions and yields a plausible-looking correlation.

Assay	L	Singles	Masked positions	Coverage
IF1_ECOLI_Kelsic_2016	72	1,367	72	100%
BLAT_ECOLX_Stiffler_2015	286	4,996	263	92%
HSP82_YEAST_Flynn_2019	709	13,194	707	100%

3.5 Full benchmark data

The three-assay analysis is superseded by a run over the complete ProteinGym substitution benchmark. All 217 assays were extracted from the **ProteinGym_v1** parquet release and every one passed wild-type validation across 2,465,767 variants. The 201 assays fitting ESM-2’s 1022-residue context were scored: 2,413,913 variants over 40,775 masked positions, spanning $L = 37$ to 934 and five function categories.

The masked-position count, not the variant count, sets the cost: masked-marginals scoring performs one forward pass per mutated position, so 2.41M variants are nearly free on top of 40,775 forwards.

An earlier three-assay path combined per-assay CSVs from the v0.1 release with wild-type sequences from the v1.0 reference file. The two agree on the 85 assays present in both and diverge elsewhere, including renames. That is acceptable for three hand-picked assays and unsafe at benchmark scale; the v1 parquet carries `target_seq` inline, so scores and wild-type sequences come from a single release.

3.6 Statistical treatment

Ground-truth correlations are compared using a **paired bootstrap** (2000 resamples). Both the candidate configuration and the fp32 reference are scored on the *same* resample at each iteration, so shared noise cancels and the residual is attributable to the configuration. The pairing is not a refinement: for `int4_wo` the paired interval is $8.4\times$ narrower than the unpaired one, because per-assay ρ spans 0.0–0.9 while the configuration effect is of order 0.006. Unpaired, every configuration reads as noise.

The *resampling unit* differs by scale of claim. For a single assay the unit is the variant, which establishes that a shift is not an artifact of the variant set. For the full benchmark the unit is the assay, since the claim becomes one about assays in general. Because the 201 assays cover only 173 distinct proteins, benchmark-level intervals use a **cluster bootstrap over proteins**: whole proteins are drawn and all their assays retained, so correlated replicates cannot count as independent

evidence. Clustered and unclustered intervals agreed to the fourth decimal, which is reported as a measured result rather than assumed.

Intervals are percentile bootstrap [5]. They were checked against a normal approximation ($\bar{d} \pm 1.96 \text{ SE}$) and agree to four decimals on every configuration, confirming the underlying distribution is near-symmetric and the percentile method is not distorting the interval. The procedure was validated on synthetic data with a known answer: a negligible perturbation was correctly classified as noise, a genuine improvement as real.

Algorithm 2 states the benchmark-level test. The pairing is line 5: one resample index is applied to the candidate and to the reference *together*. Without it the interval widens by $8.4\times$ for `int4_wo`, because between-assay variance in ρ (spanning 0.0 to 0.9) is two orders of magnitude larger than the configuration effect (≈ 0.006), and every configuration is then reported as noise.

Algorithm 2 Paired cluster bootstrap for $\Delta(\text{mean } \rho)$ against the fp32 reference. Setting $\mathcal{G}$ to singletons recovers the ordinary assay-level bootstrap; we report both.

Require: per-assay correlations $\rho_a^c, \rho_a^{\text{ref}}$ for assays $a = 1:n$; protein clusters $\mathcal{G}$ partitioning the assays; $R = 2000$

Ensure: point estimate Δ , 95% interval, two-sided p

- 1: $\Delta \leftarrow \frac{1}{n} \sum_a \rho_a^c - \frac{1}{n} \sum_a \rho_a^{\text{ref}}$
 - 2: **for** $r = 1$ **to** R **do**
 - 3: draw $|\mathcal{G}|$ clusters from $\mathcal{G}$ with replacement
 - 4: $S_r \leftarrow$ concatenation of all assays in the drawn clusters $\triangleright |S_r|$ varies between replicates
 - 5: $d_r \leftarrow \frac{1}{|S_r|} \sum_{a \in S_r} (\rho_a^c - \rho_a^{\text{ref}})$ $\triangleright$ **same** S_r indexes both arms
 - 6: **end for**
 - 7: $[\ell, u] \leftarrow$ 2.5th and 97.5th percentiles of $\{d_r\}$
 - 8: $p \leftarrow 2 \min(\frac{1}{R} |\{d_r \leq 0\}|, \frac{1}{R} |\{d_r \geq 0\}|)$ $\triangleright$ floor $2/R = 0.001$; report smaller values as $p < 0.002$
 - 9: **return** $\Delta, [\ell, u], p$; call the effect real iff $0 \notin [\ell, u]$
-

3.7 Hardware and software environment

Table 3: Hardware and software environment. The glibc 2.17 ceiling is what forces torch 2.6.0 and rules out the prebuilt bitsandbytes GPU wheels, so NF4 was unavailable and is absent from the configuration list.

Component	Detail
GPU (primary)	NVIDIA H200, 143 GB, sm_90 (FP8 tensor cores)
GPU (secondary)	NVIDIA A100, 80 GB, sm_80 (no FP8)
Driver	550.54.14
PyTorch	2.6.0+cu124
Attention	SDPA
Compile mode	<code>max-autotune-no-cudagraphs</code>

The cluster runs CentOS 7.6 with glibc 2.17 on every node, including the H200 nodes, which forced several choices. PyTorch moved to `manylinux_2.28` at version 2.7, making **2.6.0 the newest installable release**. `bitsandbytes` GPU wheels require a newer glibc, so only the CPU-only 0.42.0 installs and its NF4 path is unavailable — 4-bit quantization therefore goes through `torchao`. The system GCC (4.8.5) cannot build Triton’s helper module, so `torch.compile` fails until CC/CXX are pointed at GCC 12.1.0. Finally, `$HOME` and `/project` are at quota, and the Triton, Inductor,

HuggingFace, and pip caches all default there; all were relocated to `/scratch`. `torchao` nonetheless works because its INT8/FP8 paths dispatch to `torch._int_mm` and `torch._scaled_mm`, which live inside PyTorch and support sm_90.

4 Results

4.1 Compile mode determines whether W8A8 is viable at all

The single most consequential configuration choice is not the quantization scheme but the `torch.compile` mode. Measured on a single ESM2-3B-shaped feed-forward block, $2560 \rightarrow 10240 \rightarrow 2560$, at batch 32×512 on an A100:

Table 4: Only `max-autotune` allows Inductor to autotune the INT8 Triton matmuls, which is where the entire W8A8 benefit resides.

Variant	Time (ms)
bf16 eager	7.91
bf16 compiled	7.91
int8_dyn eager	33.35
int8_dyn compiled, <code>dynamic=True</code>	11.73
int8_dyn compiled, default mode	11.48
int8_dyn mode="max-autotune"	5.43

Measured in eager mode, W8A8 appears $4.2\times$ *slower* than bf16 and would be discarded. Measured under `max-autotune` it is $1.46\times$ faster. `dynamic=True` was avoided because it emits shape-generic kernels that surrender most of the gain; static shapes are kept manageable by padding to a fixed multiple of 128. `torch._dynamo.explain` confirms **zero graph breaks** for both bf16 and int8_dyn, establishing this as a kernel-selection effect rather than a Dynamo fallback.

4.2 Bulk embedding extraction

Table 5: ESM2-3B, H200, compiled with `max-autotune-no-cudagraphs`, SDPA attention, 64 sequences of length 512–1024 under a 16,384-token budget. Reproduced across four independent jobs: bf16 mean 56,642 residues/s, range 55,941–57,422 ($\pm 1.3\%$); fp32 mean 6,742, range 6,687–6,814.

Config	Weights (GB)	Peak (GB)	Residues/s	vs bf16	Emb. cos	Logit KL	DMS ρ_{fp32}
fp32 (ref)	10.72	25.37	6,750	$0.12\times$	—	—	—
bf16	5.29	12.63	57,057	$1.00\times$	0.999949	$1.17e-04$	0.9998
int8_wo	2.66	8.48	54,816	$0.96\times$	0.999959	$1.05e-04$	0.9998
int8_dyn	2.65	7.07	59,996	$1.05\times$	0.989592	$1.40e-02$	0.9928
fp8_dyn	2.65	7.24	65,376	$1.15\times$	0.999920	$4.60e-04$	0.9988
int4_wo	1.61	7.69	3,813	$0.07\times$	0.998904	$1.89e-03$	0.9975

Three observations. First, the **fp32 $\rightarrow$ bf16 step is worth** $8.4\times$ — larger than everything quantization adds on top of it, and the single biggest lever available. Second, **fp8_dyn** dominates **int8_dyn**: it is faster ($1.15\times$ vs $1.05\times$) *and* dramatically more accurate (embedding cosine 0.999920 vs 0.989592; logit KL 4.6×10^{-4} vs 1.4×10^{-2}), surrendering only 0.17 GB of peak memory. FP8’s E4M3 format has far greater dynamic range than INT8 at the same bit width, which matters because

transformer activations contain outliers. Third, `int4_wo` is a trap for inference: $0.07\times$ throughput and *worse* peak memory (7.69 GB) than `fp8_dyn` (7.24 GB) despite 40% smaller weights, because the dequantization path allocates transient buffers. Its value lies elsewhere, as a QLoRA base.

Separately, length-bucketed batching reduced padding waste from 25.0% (naive fixed-size batching) to **0.0%**. This is free and costs no accuracy — a larger effect than several of the quantization differences in the table.

4.3 DMS scoring: the opposite conclusion

Table 6: DMS scoring throughput, ESM2-3B, uncompiled. `bf16` wins at every problem size; quantization is pure overhead in this regime.

Config	Variants/s		
	IF1 ($L=72$)	BLAT ($L=286$)	HSP82 ($L=709$)
<code>bf16</code>	6,979	3,573	1,357
<code>int8_wo</code>	4,830	2,511	1,066
<code>fp8_dyn</code>	1,813	1,796	1,103
<code>fp32</code>	1,460	447	187
<code>int8_dyn</code>	410	382	299
<code>int4_wo</code>	1,045	282	112

`bf16` is fastest at every assay size, and `int8_dyn` — the second-fastest configuration for bulk extraction — is up to $17\times$ slower here. The masked batch ($n_{\text{positions}} \times L$) is too small to contain a matrix multiply large enough for low precision to pay for its quantize/dequantize overhead.

Notably, `fp8_dyn` closes on `bf16` as the assay grows: $0.26\times$, $0.50\times$, $0.81\times$ of `bf16` throughput at $L = 72, 286, 709$ respectively. Its overhead is fixed per call while real work scales, so at $L = 709$ it costs 20% throughput for 43% less memory — a defensible trade that is indefensible at $L = 72$.

4.4 Memory during DMS scoring

Table 7: Peak device memory during DMS scoring. Weights account for approximately 90% of peak even at $L = 709$ (`bf16`: 5.29 GB weights against 0.59 GB activations).

Config	Weights (GB)	Peak $L=72$	Peak $L=286$	Peak $L=709$
<code>fp32</code>	10.72	10.87	11.19	11.85
<code>bf16</code>	5.29	5.38	5.55	5.88
<code>int8_wo</code>	2.66	2.77	2.92	3.25
<code>int8_dyn</code>	2.65	2.87	2.91	3.36
<code>fp8_dyn</code>	2.65	2.76	2.96	3.35
<code>int4_wo</code>	1.61	1.70	1.87	2.20

This is the reverse of the bulk-extraction situation. Masked-marginals scoring keeps the batch narrow, so the $O(L^2)$ attention term never dominates and weights remain roughly 90% of peak memory even at $L = 709$. Quantization therefore addresses almost all of DMS peak memory, and `int4_wo` gives the smallest footprint at every length — $2.7\times$ below `bf16`.

Table 8: Spearman correlation against experimental measurements, ESM2-3B. Δ is the change from fp32, with a 2000-sample paired bootstrap over variants. “Real” indicates the 95% confidence interval excludes zero.

Assay	Config	ρ_{expt}	Δ vs fp32	95% CI of Δ	Verdict
IF1 ($n=1367$)	fp32	0.5561	(ref)		
	bf16	0.5543	−0.0018	[−0.0027, −0.0009]	real
	int8_wo	0.5570	+0.0009	[−0.0003, +0.0021]	noise
	int8_dyn	0.5504	−0.0057	[−0.0124, +0.0013]	noise
	fp8_dyn	0.5584	+0.0023	[−0.0022, +0.0067]	noise
	int4_wo	0.5394	−0.0167	[−0.0263, −0.0074]	real
BLAT ($n=4996$)	fp32	0.5893	(ref)		
	bf16	0.5896	+0.0003	[−0.0003, +0.0009]	noise
	int8_wo	0.5913	+0.0019	[+0.0011, +0.0029]	real
	int8_dyn	0.6160	+0.0266	[+0.0219, +0.0315]	real
	fp8_dyn	0.6088	+0.0194	[+0.0161, +0.0229]	real
	int4_wo	0.6388	+0.0495	[+0.0433, +0.0561]	real
HSP82 ($n=13194$)	fp32	0.2931	(ref)		
	bf16	0.2933	+0.0002	[−0.0002, +0.0005]	noise
	int8_wo	0.2933	+0.0002	[−0.0002, +0.0007]	noise
	int8_dyn	0.2937	+0.0006	[−0.0020, +0.0028]	noise
	fp8_dyn	0.2944	+0.0012	[−0.0003, +0.0027]	noise
	int4_wo	0.2904	−0.0027	[−0.0057, +0.0002]	noise

4.5 Ground-truth accuracy on ProteinGym

Two findings emerge, and the second qualifies the first.

Quantization does not systematically degrade variant-effect accuracy. Of the 15 assay/configuration pairs, 11 shifts are upward. Most are statistically indistinguishable from zero. Strikingly, `int4_wo` — the worst configuration on every fidelity metric — produces the *best* agreement with experiment on BLAT (+0.0495, CI [+0.0433, +0.0561]). This survives the bootstrap and is not a resampling artifact.

Fidelity-vs-fp32 predicts the magnitude of the shift, not its sign. Across all 15 pairs, the correlation between fidelity loss ($1 - \rho_{\text{fp32}}$) and the absolute ground-truth shift is $r = 0.87$ (Pearson), $\rho = 0.76$ (Spearman). But `int4_wo` moves +0.0495 on BLAT and −0.0167 on IF1. The direction is a property of the assay, not of the configuration, and is not predictable from any drift-from-fp32 measurement.

The practical consequence is that `bf16` and `int8_wo` shift ground-truth agreement by at most 0.002 on any assay tested — below anything an experiment could resolve — and are therefore safe. `int4_wo` is a ± 0.05 gamble on an unknown sign.

This supersedes an earlier conclusion drawn from a synthetic 40-residue mutational scan, which ranked `int8_dyn` ($\rho_{\text{fp32}} = 0.9928$) as risky and `fp8_dyn` (0.9988) as clearly safer. Both numbers were correct; neither predicted experimental accuracy. On BLAT the “risky” configuration gained +0.027.

We note that the assay itself matters far more than any quantization decision: ρ_{expt} ranges from 0.29 (HSP82) to 0.59 (BLAT) at 3B. Choosing which assay to trust is a $2\times$ effect; choosing a quantization configuration is a 1% effect.

4.6 Model-scale consistency

The same three assays were run on ESM2-650M. The qualitative picture holds: bf16 fastest, `int8_dyn` slowest, ground-truth shifts small. The largest 650M shift was -0.0063 (`int8_dyn` on IF1) against ρ_{fp32} of 0.9954. The larger and more clearly significant shifts at 3B indicate that sensitivity to quantization grows with model size, so 650M results should not be extrapolated to 3B without checking.

An incidental observation deserves note, because it bears directly on how much any of this matters. The fp32 650M model correlates with experiment *better* than the fp32 3B model on two of the three assays:

Table 9: Ground-truth correlation by model scale on the three-assay subset. The 0.14 BLAT gap prompted the inference that model choice dominates precision choice; Section 4.7 shows it does not generalise.

Assay	650M ρ_{expt}	3B ρ_{expt}	Difference
IF1	0.5987	0.5561	-0.0426
BLAT	0.7315	0.5893	-0.1422
HSP82	0.2648	0.2931	$+0.0283$

The BLAT gap of 0.14 is roughly *three times* the largest effect any quantization configuration produced in this study. Model selection dominates precision selection by a wide margin, and the assumption that the larger checkpoint is the more accurate one does not hold uniformly here.

This inference does not survive Section 4.7. Over 201 assays the two models are statistically indistinguishable ($p = 0.35$). The BLAT gap is real for BLAT and does not generalise — the same error, made the same way, as the `int4_wo` result three subsections above.

4.7 The full benchmark: 201 assays

Every accuracy conclusion above rests on three assays. Three is enough to establish that an effect exists on a particular assay and not enough to establish anything about assays in general, because the unit that varies — the assay — is held fixed. The whole ProteinGym substitution benchmark was therefore run: all 217 assays validated, the 201 within ESM-2’s 1022-residue context scored, 2,413,913 variants over 40,775 masked positions, six configurations, two model scales. 3.1 GPU-hours at 3B.

Two changes to the statistical treatment follow from the larger sample. The resampling unit becomes the **assay** rather than the variant, since the question is now whether a configuration degrades prediction in general. And because the 201 assays cover only 173 distinct proteins — BLAT_ECOLX appears four times — the reported interval is a **cluster bootstrap over proteins**, which prevents correlated replicates from counting as independent evidence. The two intervals turned out to agree to the fourth decimal, so the clustering was immaterial here; it is reported because that had to be measured rather than assumed.

The benchmark mean is the wrong safety criterion

Every Δ in that table is below 0.007. The per-assay tail is not (Figure 1).

`int8_dyn` at 3B is indistinguishable from fp32 on the benchmark mean ($\Delta = -0.0021$, $p = 0.34$) while destroying UBR5_HUMAN_Tsuboyama_2023_1I2T, where correlation with experiment falls from 0.591 to 0.223 at $\rho_{\text{fp32}} = 0.393$. A configuration can pass the benchmark average and be unusable on the one protein a given project cares about. This is the practical result of the whole study:

Table 10: Mean Spearman against experiment over 201 assays. Bold marks intervals excluding zero under the protein-clustered paired bootstrap. Note that `int4_wo` is significant in *opposite directions* on the two model scales. This is not driven by a handful of outliers: `int4_wo` is worse than `fp32` on 143 of 201 assays at 650M and on only 91 of 201 at 3B.

Config	ESM2-3B			ESM2-650M		
	$\bar{\rho}$	Δ	95% CI	$\bar{\rho}$	Δ	95% CI
<code>fp32</code>	0.4398	—	(ref)	0.4459	—	(ref)
<code>bf16</code>	0.4399	+0.0002	[−0.0001, +0.0004]	0.4460	+0.0000	[−0.0003, +0.0004]
<code>int8_wo</code>	0.4399	+0.0001	[−0.0002, +0.0005]	0.4457	−0.0002	[−0.0006, +0.0002]
<code>fp8_dyn</code>	0.4403	+0.0005	[−0.0007, +0.0017]	0.4452	−0.0008	[−0.0020, +0.0005]
<code>int8_dyn</code>	0.4376	−0.0021	[−0.0072, +0.0017]	0.4440	−0.0020	[−0.0032, −0.0009]
<code>int4_wo</code>	0.4435	+0.0037	[+0.0007, +0.0068]	0.4395	−0.0064	[−0.0105, −0.0018]

Table 11: Worst single-assay change in ground-truth correlation, and count of assays moving by more than 0.05, out of 201.

Config	ESM2-3B		ESM2-650M	
	worst Δ	$ \Delta > 0.05$	worst Δ	$ \Delta > 0.05$
<code>bf16</code>	−0.0072	0	−0.0040	0
<code>int8_wo</code>	−0.0106	0	−0.0120	0
<code>fp8_dyn</code>	−0.0344	0	−0.0316	0
<code>int8_dyn</code>	−0.3688	6	−0.0411	0
<code>int4_wo</code>	−0.0556	5	−0.2024	6

benchmark-mean significance testing answers a question about populations, and deployment is a question about a specific instance.

What the larger sample changed

The three-assay verdict was backwards. Section 4.5 concluded that quantization does not systematically degrade variant-effect accuracy, on the evidence that 11 of 15 shifts were upward and that `int4_wo` gained +0.0495 on BLAT. At benchmark scale `int4_wo` is significant in opposite directions on the two model scales. Both measurements were correct; the generalisation from three assays was not.

Fidelity predicts magnitude, not direction. Across 1005 assay/configuration pairs, $r(1 - \rho_{\text{fp32}}, |\Delta|) = 0.81$ at 3B and 0.74 at 650M, confirming the 0.87 estimated from 15 pairs. The *signed* correlation is −0.58 at 3B and +0.10 at 650M — inconsistent in sign between model scales, and therefore useless for prediction. ρ_{fp32} remains the correct cheap screen, bounding how far an answer can move while saying nothing about which way (Figure 2).

Damage concentrates where the model is weakest. `int4_wo`’s cost is monotone in MSA depth at *both* model scales: low < medium < high, with −0.0174 on low-depth assays against −0.0042 on high-depth at 650M. Quantization does most harm where there is least evolutionary signal to lose. This is the only stratification we find that replicates across scale — the breakdown by function category does not, and neither is a distinction three assays can support. Table 12 gives both.




Figure 1: Per-assay change in ground-truth correlation for each configuration, across all 201 ProteinGym assays and both model scales. Black ticks mark the benchmark mean, which is within 0.007 of zero everywhere. **bf16** and **int8_wo** are tight at every assay; **int8_dyn** and **int4_wo** have means that are equally unremarkable and tails that reach -0.37 and -0.20 . Selection on the mean cannot distinguish these cases; selection on the tail separates them immediately.

Speed and memory over the whole benchmark

The DMS recommendation from three assays survives and strengthens: **bf16** eager, $6.2\times$ faster than **fp32** end-to-end, with ground-truth agreement never moving more than 0.018 across 402 assay/model pairs. **int8_wo** is the memory fallback at 69% of **bf16** speed and maximum drift 0.016. **int8_dyn** and **int4_wo** are not trade-offs for this workload — they are slower than **bf16** *and* carry a catastrophic tail.

5 Reproducibility Pitfalls in Quantization Benchmarking

We document four measurement defects encountered during this study. They are included in the main text, unusually for a paper, because they bear directly on its thesis: three of the four produced clean, monotonic, plausible result tables while measuring the wrong quantity, and each would have supported a confident and wrong recommendation. A reader who accepts our argument that benchmark averages can mislead should also want to know how the underlying measurements can mislead. All four are properties of the standard tooling rather than of this codebase, and we expect them to recur in any comparable study.

5.1 Compile cost inside the timed region (three occurrences)

The DMS throughput column initially ranked **fp32** fastest and **fp8_dyn** slowest — an exact inversion of the bulk-throughput column. Two successive fixes failed. Direct instrumentation finally established the cause. On 650M/A100 with a 40-residue wild type over 11 masked positions:

DMS scoring uses a much smaller input shape than the bulk-throughput batches, and under **max-autotune** each new shape costs roughly 15s to autotune. At this batch size that cost recurs




Figure 2: Drift from fp32 against change in ground-truth correlation, over 1005 assay/configuration pairs per model. *Left*: infidelity bounds the magnitude of the change — the envelope widens monotonically, and nothing far from fp32 is reliably close to it on the ground truth. *Right*: the same data, signed. The distribution is a symmetric funnel rather than a trend, and the correlation reverses sign between model scales (-0.58 at 3B, $+0.10$ at 650M). A small drift is therefore a valid safety certificate, while a large drift is not evidence of a worse model.

rather than amortizing. The timed region was therefore measuring *autotune cost*, and ranking configurations by their number of Triton kernel candidates — fp32 had the fewest and appeared fastest, fp8_dyn had 98 and appeared slowest. The fix was to score DMS against the uncompiled module.

This has a consequence beyond the benchmark: for short wild-type sequences with few mutated positions, **max-autotune** is actively counterproductive in production, not merely in measurement.

The same class of defect appeared earlier in the bulk-throughput path, where warming up on only a prefix of the batches left recompilation for unseen length buckets inside the timed region.

5.2 Peak memory attributed to the wrong model

Peak memory during DMS scoring was reported as follows, with the activation overhead implied by subtracting weights:

Each configuration’s peak included the *previous* configuration’s weights, producing the visible absurdity of bf16 requiring more peak memory than fp32. The cleanup helper deleted only its own local reference to the model; the caller retained a binding. Because `nn.Module` graphs contain reference cycles, dropping the last reference does not free them by reference counting alone — they persist as cyclic garbage until a collection runs. Inserting an explicit `gc.collect()` before each peak-memory reset corrected it; overhead became a consistent 0.04–0.07 GB.

The bulk-throughput memory column was verified unaffected: bf16’s measured overhead there was +7.34 GB, which cannot contain fp32’s 10.72 GB of weights.

5.3 Silent argument leakage in the job script

A batch script used `shift 3 || true`. When fewer than three positional arguments are supplied, `shift 3` fails, `|| true` suppresses the failure, and the arguments remain in "\$@" — where they

Table 12: Mean ρ against experiment, stratified by ProteinGym function category and by MSA depth. **bf16**, **int8_wo** and **fp8_dyn** track **fp32** in every stratum. **int4_wo** does not. Its deviations by *function category* do not transfer across model scale — Binding is its worst stratum at 650M (-0.018) and its second-best at 3B ($+0.005$) — so category-level behaviour is no more predictable than the overall sign. Its ordering by *MSA depth* does transfer, and is monotone at both scales: low $<$ medium $<$ high ($-0.017 < -0.005 < -0.004$ at 650M, $-0.001 < +0.001 < +0.010$ at 3B). Quantization costs most where the model has least evolutionary signal to work with, and that is the one stratification here that replicates.

Stratum	fp32	bf16	int8_wo	fp8_dyn	int8_dyn	int4_wo
<i>ESM2-3B</i>						
Activity	0.4269	0.4271	0.4274	0.4267	0.4270	0.4287
Binding	0.3392	0.3396	0.3400	0.3440	0.3405	0.3441
Expression	0.4288	0.4288	0.4293	0.4294	0.4356	0.4270
OrganismalFitness	0.4006	0.4005	0.4006	0.4005	0.3990	0.4003
Stability	0.5091	0.5095	0.5089	0.5097	0.5023	0.5192
MSA depth: low	0.3264	0.3265	0.3272	0.3261	0.3239	0.3253
MSA depth: high	0.4874	0.4876	0.4879	0.4894	0.4888	0.4971
<i>ESM2-650M</i>						
Activity	0.4332	0.4329	0.4327	0.4315	0.4312	0.4217
Binding	0.3560	0.3571	0.3565	0.3562	0.3542	0.3376
Expression	0.4398	0.4397	0.4398	0.4382	0.4390	0.4398
OrganismalFitness	0.3960	0.3959	0.3959	0.3945	0.3951	0.3844
Stability	0.5233	0.5235	0.5230	0.5239	0.5200	0.5258
MSA depth: low	0.3180	0.3184	0.3186	0.3181	0.3177	0.3006
MSA depth: high	0.5176	0.5175	0.5171	0.5175	0.5151	0.5134

were appended to the Python invocation as stray arguments. Relying on a default for the third argument was sufficient to trigger it.

5.4 Common thread

Three of the four defects shared a signature: a cost that belonged outside the measurement leaked inside it, and the resulting ranking was plausible enough to be believed. In each case the tell was a result that inverted or contradicted a neighbouring measurement. We would emphasise that a benchmark producing a *clean, monotonic, plausible* table is not thereby correct; the memory bug produced exactly such a table for several runs.

Table 13: Full 201-assay benchmark, ESM2-3B on one H200, uncompiled. Wall-clock is summed scoring time over all assays.

Config	Wall-clock	Positions/s	Peak RAM (median)	Peak RAM (max)
fp32	45.0 min	25.6	11.10 GB	12.20 GB
bf16	7.3 min	200.6	5.50 GB	6.05 GB
int8_wo	9.2 min	139.3	2.87 GB	3.42 GB
fp8_dyn	10.2 min	89.2	2.90 GB	3.56 GB
int8_dyn	40.0 min	20.1	2.94 GB	3.51 GB
int4_wo	73.8 min	17.3	1.82 GB	2.37 GB

Table 14: Compilation cost inside the timed region. A single forward pass at the DMS input shape costs 15 s to autotune and 28 ms to execute, so a timed region containing compilation ranks configurations by their number of Triton kernel candidates rather than by their speed.

Measurement	Time
eager, one forward at shape (11, 42)	27.9 ms
eager, full <code>masked_marginals</code> pass	0.03 s
compiled, one forward at shape (11, 42)	15,138.1 ms

Table 15: The peak-memory artifact. Each configurations measured peak contains the previous configurations weights, which surfaces as bf16 apparently needing more peak memory than fp32.

Config	Weights (GB)	Peak (GB)	Implied overhead
fp32	2.49	2.56	+0.07
bf16	1.21	3.74	+2.53 $\approx$ fp32’s weights
int8_wo	0.61	1.86	+1.25 $\approx$ bf16’s weights
int8_dyn	0.61	1.26	+0.65 $\approx$ int8_wo’s weights

6 Discussion

6.1 The two workloads want opposite configurations

This is the principal practical finding.

Bulk embedding extraction is compute-bound: large batches, large GEMMs. Quantization pays. `fp8_dyn` with `max-autotune` wins on all three axes — $1.15\times$ bf16 throughput, weights $5.29 \rightarrow 2.65$ GB, peak $12.63 \rightarrow 7.24$ GB, with accuracy indistinguishable from bf16.

DMS variant-effect scoring is latency-bound: the masked batch is small, no GEMM is large enough for low precision to help, and per-linear quantize/dequantize overhead is pure loss. bf16 in eager mode is fastest at every problem size, and `max-autotune` is actively harmful.

Applying the embedding-optimized configuration to DMS scoring makes it up to an order of magnitude slower. Because both workloads commonly appear in the same pipeline, this is a live risk rather than a theoretical one. Figure 3 shows the two frontiers side by side.

6.2 On the interpretation of fidelity metrics

The conventional validation practice — quantize, measure drift against the full-precision model, accept if drift is small — is sound as a *risk bound*. Our data support the bound at scale: over 1005 assay/ configuration pairs, fidelity loss correlates with the magnitude of ground-truth change at $r = 0.81$ (3B) and 0.74 (650M), consistent with the 0.87 first estimated from 15 pairs.

It is not sound as a *quality ranking*. A configuration with larger drift is not thereby less accurate against reality; it is merely less predictable. The signed correlation is -0.58 at 3B and $+0.10$ at 650M — it does not even hold its sign between model scales.

A tempting misreading of our BLAT result is that quantization acts as beneficial regularization. We do not endorse this, and the full benchmark settles it: `int4_wo` is significantly *positive* at 3B and significantly *negative* at 650M. A configuration cannot be genuinely more accurate than the reference from which it is derived in any way one could rely on in advance.




Figure 3: Speed–memory frontiers for the two workloads, ESM2-3B on one H200; lower-right is better in both panels. *Left*: bulk embedding extraction under `max-autotune`, where `fp8_dyn` is fastest and near-smallest. *Right*: DMS scoring over the full 201-assay benchmark in eager mode, where the same configuration is $2.2\times$ slower than `bf16` and the ranking is substantially reordered. No single configuration is preferred by both.

6.3 Averages conceal the failure that matters

The benchmark-mean result is that no configuration changes mean ground-truth correlation by more than 0.007. Read alone, that licenses any configuration on the list. It should not.

`int8_dyn` at 3B is statistically indistinguishable from `fp32` on the mean ($p = 0.34$) and takes one assay from $\rho = 0.591$ to 0.223. A practitioner does not deploy against the mean of 201 assays; they deploy against one protein. The benchmark mean answers whether a configuration is safe *on average over targets*, which is a claim about a population, while the operative question is whether it is safe on a specific target.

The practical consequence is that the tail statistics — worst-case single-assay drift, and the count of assays exceeding a tolerance — are the numbers to select on, and they separate the configurations far more sharply than the means do. On that criterion `bf16` and `int8_wo` are safe (no assay beyond 0.018 across 402 assay/model pairs), `fp8_dyn` is nearly so (0.034), and `int8_dyn` and `int4_wo` are not.

6.4 Recommendations

Two non-quantization measures outweigh most of the above. Moving from `fp32` to `bf16` is worth $8.4\times$ on bulk extraction and $6.2\times$ over the full DMS benchmark; length-bucketed batching eliminates 25% of wasted computation. Both should be applied before any quantization is considered.

A third, proposed in an earlier draft of this report, does not survive the full benchmark: the 650M checkpoint beat 3B by 0.14 Spearman on BLAT, which suggested model selection dominated precision selection. Over 201 assays the two are indistinguishable (0.4459 vs 0.4398, 95% CI $[-0.0061, +0.0185]$, $p = 0.35$, with 650M ahead on 110 of 201). What remains true is narrower and still useful: the *uncertainty* on model choice (± 0.012) exceeds every quantization effect measured here (≤ 0.007). Precision is the settled variable; checkpoint choice is not, and is worth benchmarking

Table 16: Configuration recommendations by situation. The first two rows differ because the two workloads sit in different bottleneck regimes, not because of any accuracy difference.

Situation	Recommendation
Bulk embedding extraction (H200)	<code>fp8_dyn + max-autotune-no-cudagraphs</code>
Bulk embedding extraction (A100)	<code>int8_wo</code> ; <code>int8_dyn</code> costs real accuracy
DMS / variant effect scoring	<code>bf16</code> , eager, no compile
DMS under memory pressure	<code>int8_wo</code> (2.87 vs 5.50 GB; 69% of bf16 speed)
Minimum memory footprint	<code>int4_wo</code> , only if a -0.20 single-assay collapse is acceptable; validate on your target
QLoRA fine-tuning base	<code>int4_wo</code> (1.61 GB weights)
V100 / sm_70 hardware	Quantization buys memory only, never speed

on a representative assay.

6.5 A selection protocol

Our results imply a procedure rather than a ranking, because the quantity that separates configurations — worst-case single-assay drift — is cheap to measure on a target and cannot be inferred from published means. Algorithm 3 states it. The essential point is line 4: the screen runs against the fp32 model on the user’s own sequences and needs no experimental labels, so it costs one extra scoring pass and no wet-lab work.

Algorithm 3 Choosing a precision configuration for a variant-effect workload. The screen requires no ground-truth labels, because Section 6.2 establishes that drift from fp32 is a valid bound on how far the answer can move even though it does not indicate direction.

Require: target protein(s) S , tolerance τ on Spearman drift, memory budget m

- 1: score S under `bf16`; **if** it fits in m , **return** `bf16` $\triangleright$ fastest at every problem size we measured
 - 2: $\mathcal{C} \leftarrow \langle \text{int8_wo}, \text{fp8_dyn} \rangle$ $\triangleright$ ordered by measured worst-case drift, not by mean
 - 3: **for** $c \in \mathcal{C}$ with weight footprint $\leq m$ **do**
 - 4: $\rho_c \leftarrow \text{Spearman}(\sigma_c(V_S), \sigma_{\text{fp32}}(V_S))$ $\triangleright$ one scoring pass; no labels needed
 - 5: **if** $1 - \rho_c \leq \tau$ **then return** c
 - 6: **end for**
 - 7: **if** m still unmet: use `int4_wo` **only** after confirming $1 - \rho_c \leq \tau$ on *this* target $\triangleright$ worst observed single-assay collapse: -0.20
 - 8: **else** reduce m by reducing model scale before reducing precision
-

We deliberately do not recommend selecting on published benchmark means. On the 201-assay mean, `int8_dyn` and `bf16` are separated by 0.002; on worst-case single-assay behaviour they are separated by 0.36. The mean is the statistic most likely to be quoted and the least useful for this decision.

6.6 Computational cost

The full benchmark costs 3.1 GPU-hours at 3B and 0.9 at 650M on a single H200 for all six configurations — 40,775 masked forward passes per configuration, six configurations, two scales. This is small enough that evaluating on the complete benchmark rather than a subset is a matter of choosing to, not of affording to. The three-assay evaluation this study began with saved roughly two GPU-hours and produced two conclusions that the full benchmark reversed.

7 Limitations

- **We still cannot predict which assays behave which way.** 201 assays establish the population behaviour of each configuration and the size of its tail. They do not identify in advance which target will be the one that collapses. The only reliable procedure remains scoring a candidate configuration against fp32 on the actual target.
- **The 16 longest assays are excluded.** Assays beyond 1022 residues exceed ESM-2’s context and were skipped rather than truncated, since a truncation window is a scoring-protocol change that would confound the comparison. Configuration behaviour above $L = 934$ is therefore unmeasured.
- **Multiple comparisons are uncorrected.** Five configurations are tested against fp32 per model. Both significant results survive a Bonferroni threshold of $0.05/5 = 0.01$, but `int4_wo` at 650M sits exactly on it ($p = 0.010$).
- **Bootstrap resolution.** At 2000 resamples the smallest reportable two-sided p is 0.001; values at that floor should be read as $p < 0.002$ rather than as point estimates.
- **Low-signal assays resolve nothing individually.** Where $\rho_{\text{expt}} \approx 0.3$ the model’s own error dominates and configuration differences fall below the noise floor; such assays contribute to the benchmark mean but carry little information alone.
- **Sequence-length coverage.** Bulk-throughput measurements used sequences of 512–1024 residues. Beyond $L \approx 2000$ the $O(L^2)$ attention term dominates peak memory and quantization’s share of the memory benefit shrinks correspondingly. A representative FASTA from the target workload would settle this.
- **Single random seed** for quantization; no ensemble over calibration variation.
- **The QLoRA track is unbuilt.** `int4_wo` at 1.61 GB is the natural base, and fine-tuning is the one setting where quantization buys capability rather than efficiency — full 3B fine-tuning requires roughly 45 GB of optimizer state alone.

8 Conclusions

What to run. Two workloads, two different answers, and no compromise configuration.

- **Bulk embedding extraction: fp8_dyn with max-autotune.** Best on all three axes simultaneously — $1.15\times$ bf16 throughput, peak memory $12.63 \rightarrow 7.24$ GB, accuracy indistinguishable from bf16. This requires FP8 hardware (`sm_90`); on A100, `int8_wo`.
- **Variant-effect scoring: bf16, eager, uncompiled.** Fastest at every problem size from $L = 37$ to 934, and $6.2\times$ faster than fp32 over the full benchmark. Quantization here is pure overhead and compilation is actively harmful. Use `int8_wo` if memory forces it (2.87 vs 5.50 GB, 69% of bf16 speed).
- **Avoid `int8_dyn` and `int4_wo` for variant-effect work.** They are not trade-offs: both are *slower* than bf16 in this regime and carry single-assay collapses of -0.37 and -0.20 .

What the evidence says about evidence. Three findings generalise past ESM-2, and each contradicts a common practice.

1. **Select on worst-case drift, not on benchmark means.** Across 201 assays no configuration moves mean correlation by more than 0.007 — the statistic most likely to be published, and the one least able to distinguish these configurations. `int8_dyn` passes the mean test at 3B ($p = 0.34$) and takes one assay from $\rho = 0.591$ to 0.223. On means, `int8_dyn` and `bf16` differ by 0.002; on worst case, by 0.36 — a factor of 180.
2. **Drift from fp32 is a valid safety certificate and an invalid quality ranking.** Over 1005 assay/configuration pairs it predicts the *magnitude* of ground-truth change ($r = 0.81$) but not its *direction*, whose correlation reverses sign between model scales (-0.58 at 3B, $+0.10$ at 650M). So a configuration close to fp32 is safe, while a configuration far from fp32 is merely unpredictable — not worse. `bf16` and `int8_wo` never exceed 0.018 Spearman on any of 402 assay/model pairs and are safe on that basis alone; this is cheap to verify on a target and needs no experimental labels (Algorithm 3).
3. **Evidence scale is not a matter of rigour but of conclusions.** This study reached three successive verdicts. A synthetic mutational scan ranked configurations confidently and wrongly. Three real ProteinGym assays overturned it and produced two further generalisations — that quantization does not systematically degrade accuracy, and that checkpoint choice dominates precision choice — both of which the full 201-assay benchmark reversed. Every intermediate conclusion was statistically sound on its own data and looked well-supported. The full benchmark costs 3.1 GPU-hours; the three-assay version saved two and cost two wrong conclusions.

The general claim. A benchmark mean answers whether a configuration is safe *on average over targets*. Deployment asks whether it is safe on *one* target. These questions have different answers here, and the gap between them is not a detail — it is the difference between 0.002 and 0.36. Wherever a benchmark aggregates over heterogeneous instances and the user faces one instance, we expect the same gap, and we would report the tail alongside the mean as a matter of course.

9 Data and Code Availability

All assay data is redistributed from ProteinGym [11], release v1, obtained from the OATML-Markslab/ProteinGym_v1 dataset repository. No new experimental measurements were produced.

The following are released with this work:

- Benchmark harness, scoring driver, and analysis code, including the extraction script that validates every variant’s wild-type residue against `target_seq` before writing.
- **Per-variant predictions** for all 12 model/configuration combinations: 2,413,913 scores each, in the canonical assay row order, so that any alternative statistic can be recomputed without access to a GPU.
- Per-assay statistics (1206 rows per model: accuracy, wall-clock, peak memory, and fidelity for every assay/configuration pair) and the aggregated summaries reported in Section 4.

[TODO before submission: insert the public repository or archive DOI here. arXiv permits a URL in this section; a Zenodo DOI is preferable for the per-variant score files, which total roughly 90 MB compressed.]

References

- [1] J. Ansel et al. PyTorch 2: Faster machine learning through dynamic Python bytecode transformation and graph compilation. In *ASPLOS*, 2024.
- [2] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *NAACL-HLT*, 2019.
- [3] T. Dettmers, M. Lewis, Y. Belkada, and L. Zettlemoyer. LLM.int8(): 8-bit matrix multiplication for transformers at scale. In *NeurIPS*, 2022.
- [4] T. Dettmers, A. Pagnoni, A. Holtzman, and L. Zettlemoyer. QLoRA: Efficient finetuning of quantized LLMs. In *NeurIPS*, 2023.
- [5] B. Efron and R. J. Tibshirani. *An Introduction to the Bootstrap*. Chapman and Hall/CRC, 1993.
- [6] E. Frantar, S. Ashkboos, T. Hoeffer, and D. Alistarh. GPTQ: Accurate post-training quantization for generative pre-trained transformers. In *ICLR*, 2023.
- [7] Z. Lin et al. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*, 379(6637):1123–1130, 2023.
- [8] J. Lin et al. AWQ: Activation-aware weight quantization for on-device LLM compression and acceleration. In *MLSys*, 2024.
- [9] J. Meier, R. Rao, R. Verkuil, J. Liu, T. Sercu, and A. Rives. Language models enable zero-shot prediction of the effects of mutations on protein function. In *NeurIPS*, 2021.
- [10] P. Micikevicius et al. FP8 formats for deep learning. *arXiv:2209.05433*, 2022.
- [11] P. Notin et al. ProteinGym: Large-scale benchmarks for protein fitness prediction and design. In *NeurIPS Datasets and Benchmarks*, 2023.
- [12] A. Rives et al. Biological structure and function emerge from scaling unsupervised learning to 250 million protein sequences. *PNAS*, 118(15):e2016239118, 2021.
- [13] P. Tillet, H. T. Kung, and D. Cox. Triton: An intermediate language and compiler for tiled neural network computations. In *MAPL*, 2019.
- [14] torchao maintainers. torchao: PyTorch-native training-to-serving model optimization. <https://github.com/pytorch/ao>, 2024.
- [15] K. Tsuboyama et al. Mega-scale experimental analysis of protein folding stability in biology and design. *Nature*, 620:434–444, 2023.

A Reproduction

```
# Environment (handles glibc/GCC/cache-quota constraints)
source env/activate.sh

# Fetch assay data; verifies variant indexing against target_seq
python src/fetch_proteingym.py IF1_ECOLI_Kelsic_2016 \
    BLAT_ECOLX_Stiffler_2015 HSP82_YEAST_Flynn_2019

# Bulk throughput + fidelity matrix (compiled)
sbatch -p H8V141_SAP112M2000_L --mem=200G slurm/run_matrix.sbatch \
    3B fp32,bf16,int8_wo,int8_dyn,fp8_dyn,int4_wo --compile

# Real-assay DMS: accuracy, speed, memory (uncompiled by design)
sbatch -p H8V141_SAP112M2000_L --mem=200G slurm/run_dms.sbatch \
    3B fp32,bf16,int8_wo,int8_dyn,fp8_dyn,int4_wo
```

```

# Significance of ground-truth shifts (single assay, variant bootstrap)
python src/analyze_dms.py results/dms_3B_*-<jobid>_scores.json

# --- Full benchmark -----
# All 217 assays; validates every variant's WT and refuses to write on
# mismatch. ~110 MB of parquet, cached in data/proteingym_v1/_raw
python src/fetch_proteingym_v1.py

# One array task per config; each loads its model once and scores all
# 201 assays. Resumable: a resubmission skips assays already written.
sbatch --array=0-5 -p H8V141_SAP112M2000_L --mem=200G \
    slurm/run_proteingym.sbatch 3B
sbatch --array=0-5 -p H8V141_SAP112M2000_L --mem=200G \
    slurm/run_proteingym.sbatch 650M

# Benchmark-wide accuracy / speed / RAM + protein-clustered bootstrap
python src/aggregate_proteingym.py --model 3B
python src/aggregate_proteingym.py --model 650M

```

Result artifacts referenced in this report:

- Throughput matrices (SLURM jobs 5393462, 5393520, 5393524, 5393554):
results/matrix_3B-<jobid>.json
- ProteinGym, 3B, with per-variant scores (job 5394386):
results/dms_3B-<assay>_5394386.json
- ProteinGym, 650M (job 5394316): results/dms_650M-<assay>_5394316.json
- Full benchmark, per-variant scores, 3B (array 5395036) and 650M (array 5395037):
results/proteingym/<model>_<config>.jsonl.gz
- Full-benchmark analysis: results/proteingym/summary-<model>.json and
summary-<model>_per_assay.csv (1206 rows: one per assay/configuration)